

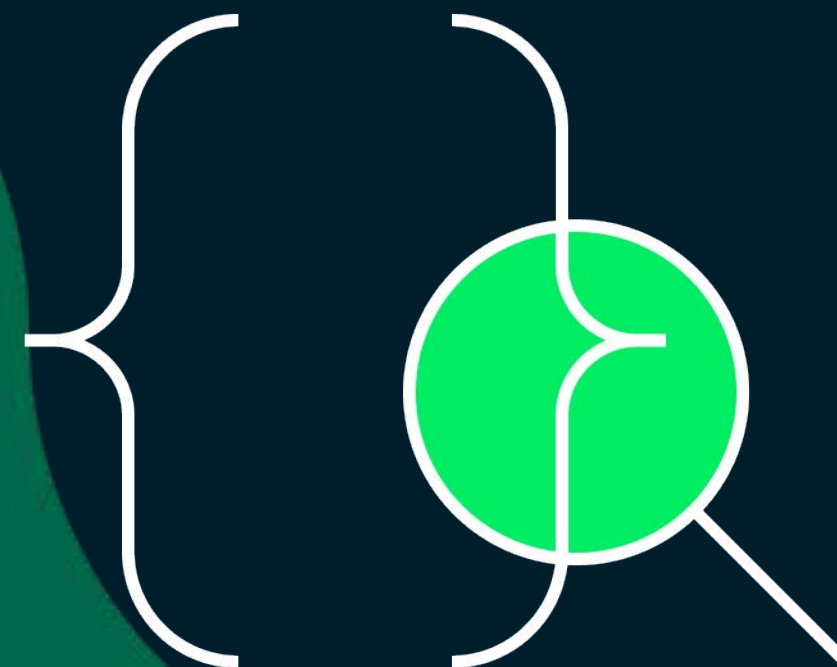
Theory & Demo

Kai Yong Lai (Vandyck)

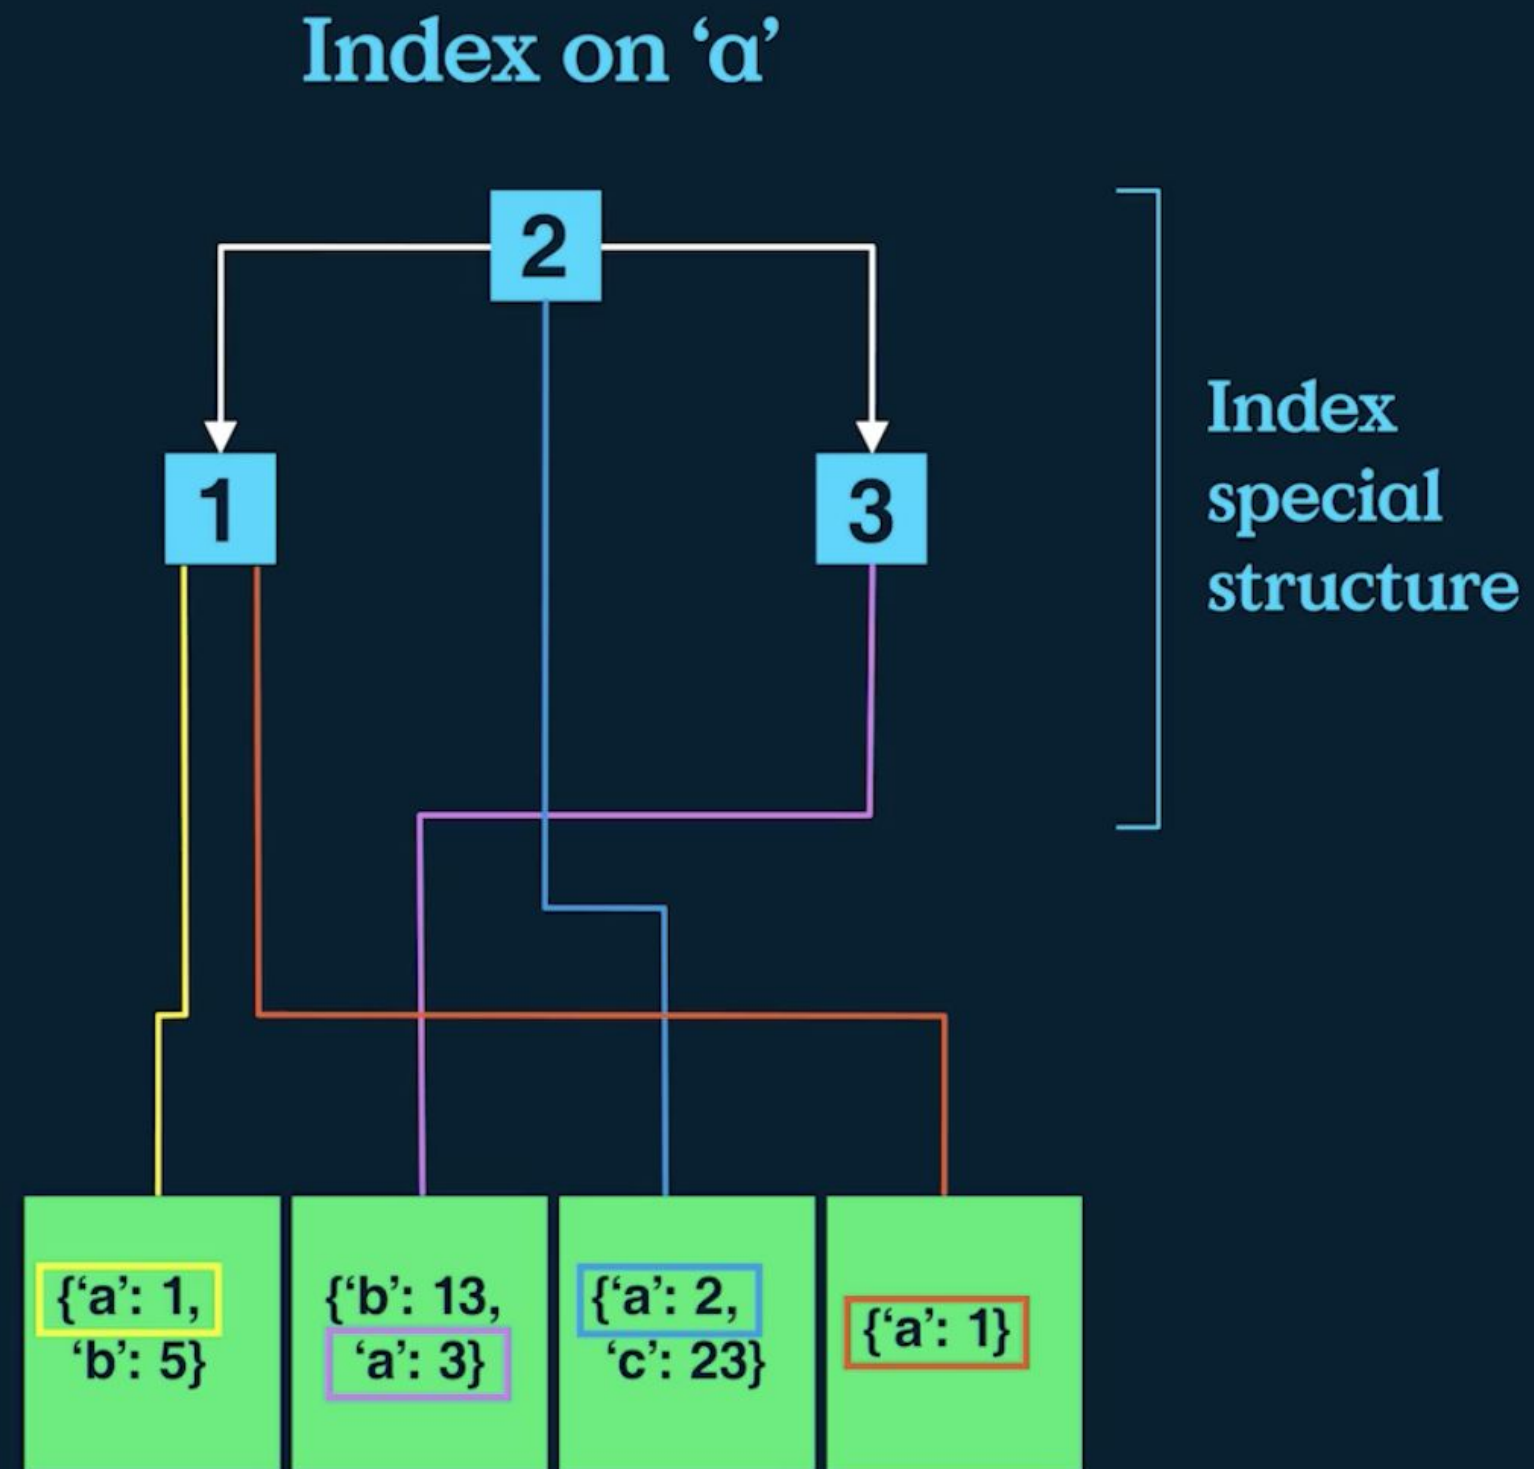
Is your index INDEXING?



Community Creator + User Group Leader



Basic of Indexing

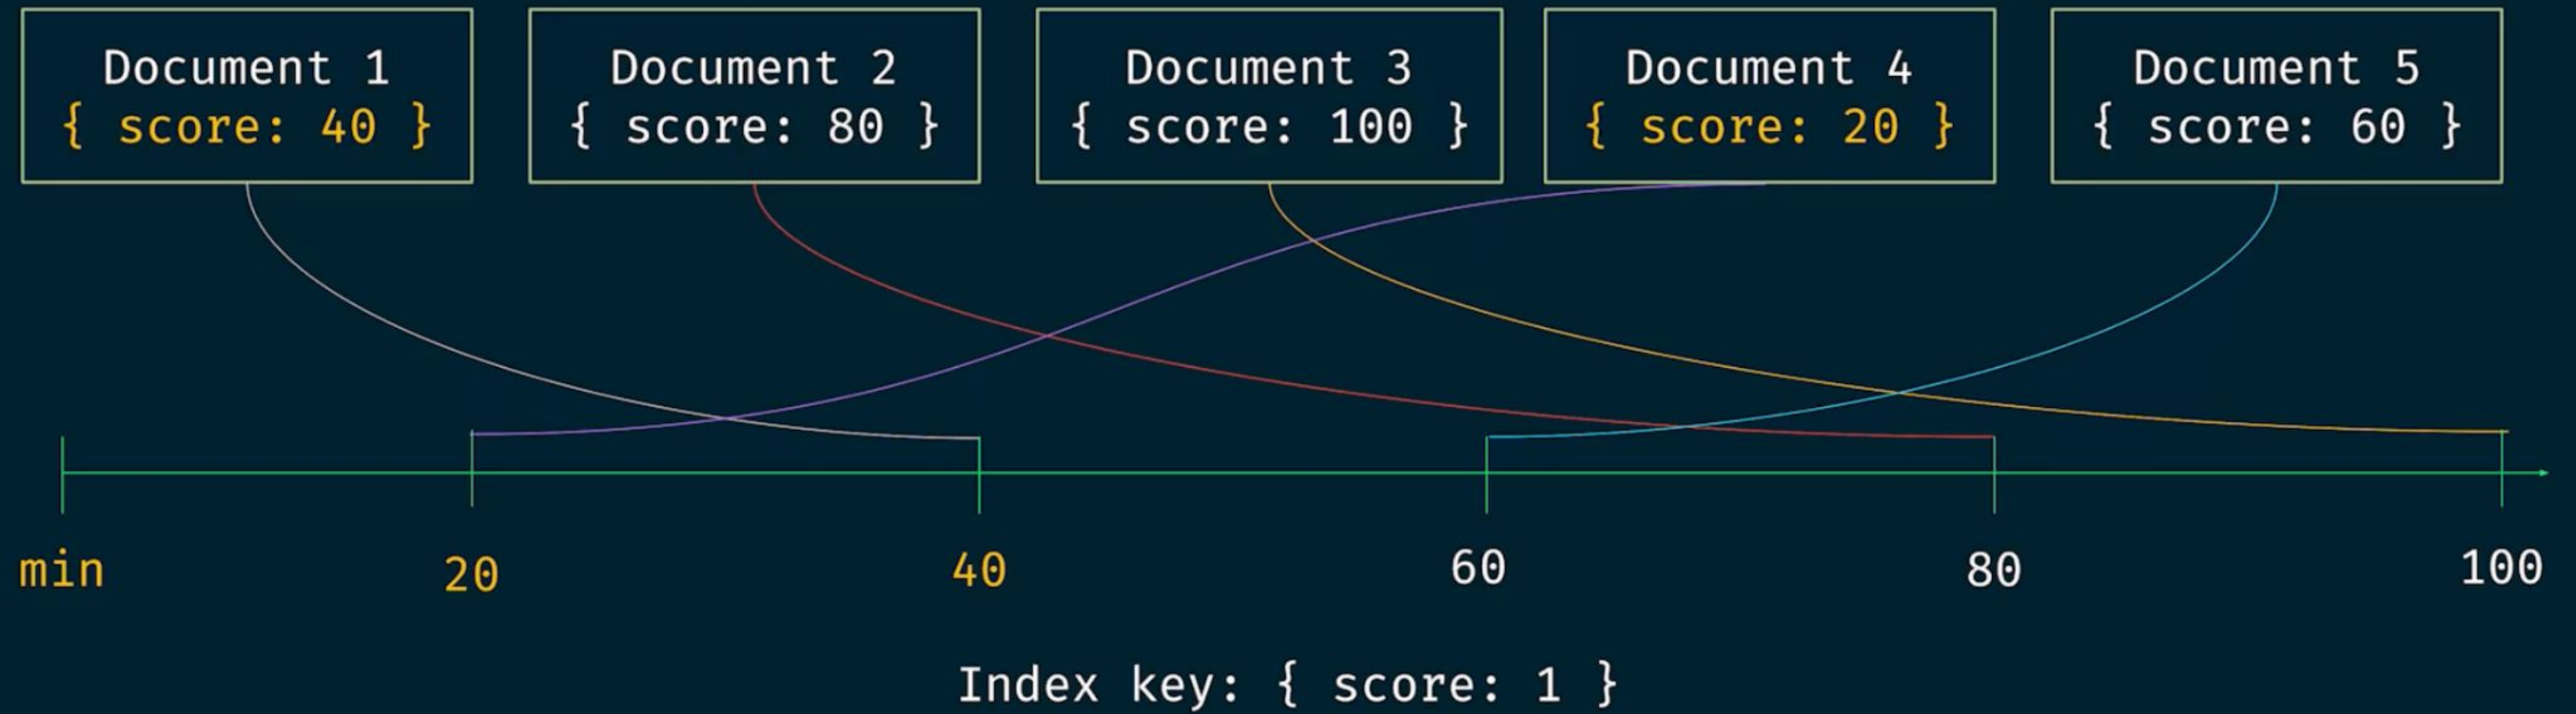


Data structure for **fast RETRIEVAL**



How

Query: { score : { \$lte: 40 } }



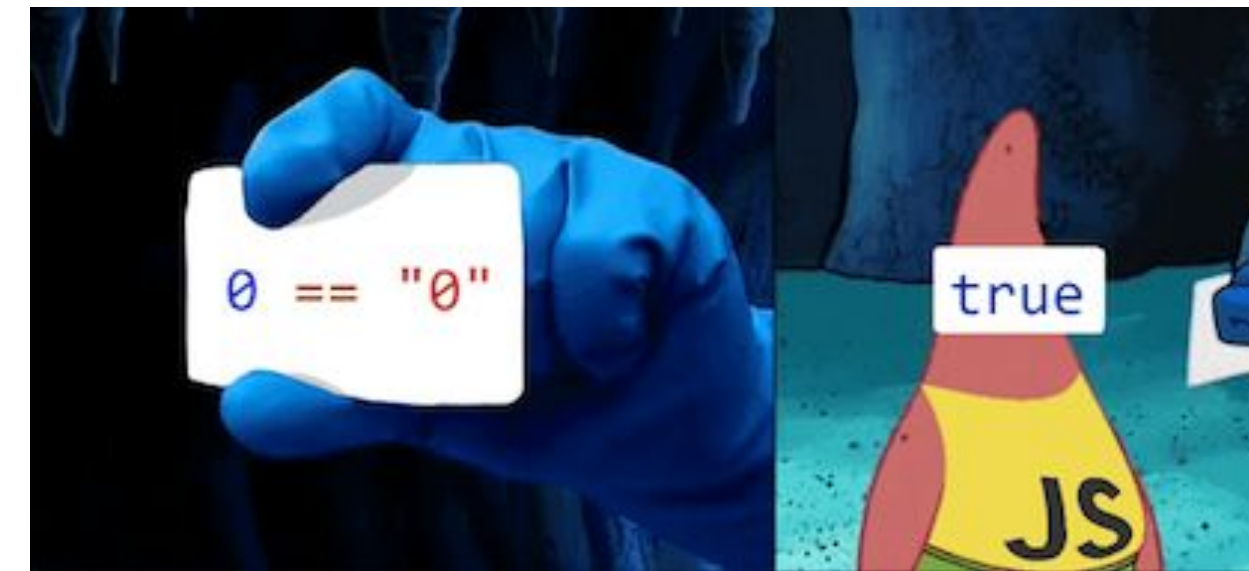
Why Index



Speed



Resource & Disk IO



Equality Operator
+ Range

Single Field



```
_id: ObjectId('5a9427648b0beeb69579e7')
name: "Mercedes Tyler"
email: "mercedes_tyler@fakegmail.com"
movie_id: ObjectId('573a1390f29313caabcd4323')
text: "Eius veritatis vero facilis quaerat fuga temporibus.
      Praesentium exped..."
date: 2002-08-18T04:56:07.000+00:00
```

Example:

```
db.comments.createIndex({
  name: 1
})
```

MultiKey



```
_id: ObjectId('59b99db4cfa9a34dcd7885b6')  
name : "Ned Stark"  
email : "sean_bean@gameofthron.es"  
password : "$2b$12$UREFwsRUoyF0CRqGNK0Lz00HM/jLhgUCNN..."
```

Example:

```
db.users.createIndex({  
  email: 1  
})
```

Compound



```
_id: ObjectId('5a9427648b0beebeb69579e7')
name: "Mercedes Tyler"
email: "mercedes_tyler@fakegmail.com"
movie_id: ObjectId('573a1390f29313caabcd4323')
text: "Eius veritatis vero facilis quaerat fuga temporibus.
      Praesentium exped..."
date: 2002-08-18T04:56:07.000+00:00
```

Example:

```
db.comments.createIndex({
  name: 1,
  email: 1
})
```

WildCard



```
_id: ObjectId('59a47286cfa9a3a73e51e72c')
theaterId : 1000
▼ location : Object
  ▼ address : Object
    street1 : "340 W Market"
    city : "Bloomington"
    state : "MN"
    zipcode : "55425"
  ▼ geo : Object
    type : "Point"
    ▼ coordinates : Array (2)
      0: -93.24565
      1: 44.85466
```

Example:

```
db.theaters.createIndex({
  "location.address.$**" : 1
})
```


Partial Index



```
_id: ObjectId('573a1390f29313caabcd5293')
plot: "Young Pauline is left a lot of money when her w
Howe..."
▶ genres: Array (1)
  runtime: 199
▶ cast: Array (4)
  num_mflix_comments: 0
  poster: "https://m.media-
amazon.com/images/M/MV5BMzgxODk1Mzk2Ml5BMl5Ban
title: "The Perils of Pauline"
fullplot: "Young Pauline is left a lot of money when h
dies. Howe..."
▶ languages: Array (1)
  released: 1914-03-23T00:00:00.000+00:00
▶ directors: Array (2)
▶ writers: Array (5)
▶ awards: Object
  lastupdated: "2015-09-12 00:01:18.647000000"
  year: 1914
▶ imdb: Object
```

Example:

```
db.movies.createIndex(
  {
    "year": 1
  },
  {
    "partialFilterExpression":
    {
      "year": {
        $gte: 1910
      }
    }
  }
)
```

Collation



```
_id: ObjectId('6778391eb518e48f1177c46e')
title: "Token Balance Validator"
description: "Create a smart contract that
             implements a basic token balance
             checking..."
difficulty: "Easy"
createdAt: 2025-01-03T19:23:10.400+00:00
testCases: "// SPDX-License-Identifier: MIT
            pragma solidity ^0.8.0
            ..."
solution: "// SPDX-License-Identifier: MIT
          pragma solidity ^0.8.0;

          contract Token..."
```

Example:

```
db.problems.createIndex(
  { "difficulty": 1 },
  {
    "collation": {
      "locale": "en",
      "strength": 2
    }
  }
)
```



Equality -> Sort -> Range

```
db.customers.find({  
  birthdate: {  
    $gte:ISODate("1977-01-01")  
  },  
  active:true  
}).sort({  
  name:1  
})
```

```
db.customers.createIndex({  
  active:1,  
  name:1,  
  birthdate: 1  
})
```

Inspection Mode

- `queryPlanner`, which details the plan selected by the `query optimizer` and lists the rejected plans.
- `executionStats`, which details the execution of the winning plan and the rejected plans.
- `serverInfo`, which provides information on the MongoDB instance.

`allPlansExecution` Mode

By default, `explain` runs in "`allPlansExecution`" verbosity mode. The following `explain` command returns the `queryPlanner` and `executionStats` for all considered plans for an `update` command:

Stage

- `COLLSCAN` for a collection scan
- `IXSCAN` for scanning index keys
- `FETCH` for retrieving documents
- `GROUP` for grouping documents
- `SHARD_MERGE` for merging results from shards
- `SHARDING_FILTER` for filtering out orphan documents from shards
- `TS_MODIFY` for modifying a time series collection
- `BATCHED_DELETE` for multiple document deletions that are batched together internally (starting in MongoDB 6.1)
- `EXPRESS` stages for a limited set of queries that can bypass regular query planning to use optimized index scan plans (*New in version 8.0.*)

`EXPRESS` stages can be one of the following:

- `EXPRESS_CLUSTERED_IXSCAN`
- `EXPRESS_DELETE`
- `EXPRESS_IXSCAN`
- `EXPRESS_UPDATE`

Code

